

Same architecture, 7.8× the data: fresh-data training

